UX Roadmap Fact-Check & Plan

Updated 2026-11-13. Summary of docs/analysis/ux_roadmap_fact_check.md (source of truth). Streamlit console at ui/, HTTP-only, API-driven.

Status: UX1-UX12 are delivered and closed (#130-#136, #145, #151-#155). The remaining tranche, UX13-UX15, is opened as #156-#158 and mapped to Project #8 Workstream M.

Assessment: 'Don't make me think'

What works well: HTTP-only client; grouped, role-aware navigation; 15-step tour and action-based Get started checklist; Query Studio templates, history, compare and result tabs; dataset evaluation with per-case isolation; trace waterfall; governance filters/export; admin safety; remediation hints; persistent workspace and shareable deep links; background evaluation jobs; Dataset Studio authoring; cross-object evidence and request diffs; keyboard-first command palette plus a WCAG 2.2 AA decision record.

Remaining friction: one generic tour for every persona; no demo mode or safe reset; no responsive/compact/high-contrast theme system; bad answers cannot be annotated or promoted into evaluation data; no saved executive/ops dashboards.

Delivered UX Workstreams (Closed)

ID	Workstream	Issue	Status	Evidence
UX1	Simplify IA and onboarding	#130	Shipped	ui/navigation.py; tests/test_ux_navigation.py
UX2	Query Studio workflow polish	#131	Shipped	ui/query_studio.py; tests/test_query_studio.py
UX3	Evaluation UX improvements	#132	Shipped	ui/evaluation.py; tests/test_ui_evaluation_ux.py
UX3b	Dataset execution follow-up	#145	Shipped	ui/evaluation.py; tests/test_ui_dataset_evaluation.py
UX4	Traces timeline/waterfall	#133	Shipped	ui/traces.py; tests/test_ui_traces_timeline.py
UX5	Governance filters and export	#134	Shipped	ui/governance.py; tests/test_ui_governance_filters.py
UX6	Admin/tenancy safety UX	#135	Shipped	ui/admin.py; tests/test_ui_admin_safety.py
UX7	Polish, a11y, empty states	#136	Shipped	ui/polish.py; tests/test_ui_polish.py
UX8	Persistent workspace/deep links	#151	Shipped	ui/workspace.py; tests/test_ui_workspace.py
UX9	Background evaluation jobs	#152	Shipped	ui/eval_jobs.py; tests/test_ui_eval_jobs.py
UX10	Dataset Studio authoring	#153	Shipped	ui/dataset_studio.py; tests/test_ui_dataset_studio.py
UX11	Cross-object evidence	#154	Shipped	ui/evidence.py; tests/test_ui_evidence.py
UX12	Command palette + WCAG 2.2 AA	#155	Shipped	ui/palette.py; ui/a11y.py; tests/test_ui_palette.py; docs/accessibility.md

Remaining UX/UI Proposals (UX13-UX15, Opened)

Opened as #156-#158. Tracking: label ux, link to umbrella #43, Project #8 Workstream M, milestone H4, HTTP-only rule plus AST/client tests for every change.

ID	Proposal	Priority	Problem	Acceptance
UX13	Responsive, compact and theme system	P2	Wide layouts break on small screens; tables/colors are inflexible; no high-contrast mode	375/768/1280 px core flows have no horizontal scroll; contrast >=4.5:1; theme che
UX14	Role-based guided journeys and demo mode	P2	One generic tour; no safe demo data or reset for personas	Each persona journey <10 min; reset is clean; tour covers new screens; no produc
UX15	Collaboration and feedback capture	P2	Bad answers cannot be annotated or promoted into evaluation data; no shareable notes	Bad-answer annotations survive reload; exports are redacted.

Recommended Execution Order

1. UX8 persistence/deep links unlocked the rest. 2. UX9 removed the biggest evaluation blocker. 3. UX12 delivered the accessibility gate. 4. UX10 made evaluation repeatable. 5. UX11 shortened debugging. 6. UX13-UX15 add presentation, persona onboarding and collaboration polish.

Bottom Line

UX1-UX7 turned the console from a functional API demo into a role-aware, guided workflow surface. UX8-UX12 made it persistent, asynchronous, evidence-linked and keyboard-accessible. The remaining tranche (#156-#158) should make it responsive, persona-specific and collaborative.